

Gaurav Arora

APPLIED SCIENTIST - 2 @ IML, AMAZON

☎ 9888552871 | ✉ contactgauravforwork@gmail.com | 🏠 www.arggaurav.com | 📱 goru001 | 📠 gaurav-arora-23593220

Education

Georgia Institute of Technology

Atlanta, Georgia

MASTERS IN COMPUTER SCIENCE WITH SPECIALIZATION IN MACHINE LEARNING

2020 - 2022

- My GPA was 3.8/4. I pursued my Masters remotely during Covid in addition to full-time job at Haptik and then at Amazon.

Punjab Engineering College

Chandigarh, India

BACHELOR OF TECHNOLOGY (B.TECH.), COMPUTER SCIENCE AND ENGINEERING

2014 - 2018

- My CGPA was 7.77/10. I had the best JEE-Mains Rank in class.

Experience

Amazon

Bengaluru

APPLIED SCIENTIST - 2

Apr 2021 - Present

- Helped in building and launching Conversational Shopping Assistant (CSA) for Amazon (India) in the pre and post LLM world.
- Built multiple NLU components (Intent, NER, Query Reformulation, Policy Engine), Safeguard service, Next Question Suggestion service, Clarification and Product QnA tools for the assistant, owning them from research to production. My work has resulted in improvements across coverage and precision of bot response upto 10%.
- Worked on improving representation of code-mixed data in Transformer models, paper for which was accepted in ACL 2023.
- Worked on improving multilingual representations in LLMs, paper for which is under review at EMNLP 2024.
- Also worked on forecasting problems using Deep Learning and built the demand forecasting system for predicting smartphones demand for Amazon India, improving the model performance by 20%.

Haptik (part of Jio Platforms Limited)

Mumbai

MACHINE LEARNING SCIENTIST - 2

July 2019 - Apr 2021

- Built the Intent Detection System for Haptik's NLU Engine, owning it from research to production.
- The performance of new system was bench-marked against 10 data-sets, including 3 open source datasets and it was 25% more accurate than their previous system, with latency under 150ms.
- We also benchmarked the new system against popular NLU platforms like Dialogflow, RASA, LUIS in the paper HINT3: Raising the bar for Intent Detection in the Wild and Haptik's system achieved similar performance as others.

Goldman Sachs

Bengaluru

SOFTWARE DEVELOPMENT ENGINEER - 1

June 2018 - July 2019

- Worked on Analytics for Desktop Assistant to bring to the fore key metrics about Product Health and User Behavior.

Researchshala

Chandigarh

CO-FOUNDER

Nov. 2016 - Mar. 2018

- Built an online platform to connect professors to research-interns, helping them with their research projects.

Publications

Towards Abstractive Knowledge Representations in LLMs

Under review @ EMNLP 2024

GAURAV ARORA, SHREYA JAIN, VAIBHAV SAXENA, SRUJANA MERUGU

- We evaluate LLMs' ability to reason using two foundational relationships: "equivalence" and "inheritance", by introducing novel tasks and evaluation benchmarks spanning six languages.
- Highlight systemic gaps in LLMs by showing that current SOTA LLMs often produce conflicting answers to the same questions across languages in 17.3-57.5% cases and also violate inheritance constraints in up to 37.2% cases.
- We propose novel "Compositional Representations" where tokens are represented as composition of equivalent tokens across languages, with resulting conflict reduction up to -4.7%.

Intent Detection in the Age of LLMs

Under review @ EMNLP 2024

GAURAV ARORA, SHREYA JAIN, SRUJANA MERUGU

(Industry Track)

- We adapt 7 SOTA LLMs using adaptive in-context learning and chain-of-thought prompting for intent detection, and compare their performance with contrastively fine-tuned sentence transformer (SetFit) models to highlight prediction quality and latency tradeoff.
- We propose a hybrid system using uncertainty based routing strategy to combine the two approaches, achieving the best of both worlds (i.e. within 2% of native LLM accuracy with 50% less latency).
- Introduce a two-step approach utilizing internal LLM representations, demonstrating empirical gains in OOS detection accuracy and F1-score by >5% for the Mistral-7B model.

CoMix: Guide transformers to code-mix using POS structure and phonetics

Accepted @ ACL 2023 (Findings)

GAURAV ARORA, SRUJANA MERUGU, VIVEK SEMBIUM

- We propose CoMix, a pre-training approach to improve representation of code-mixed data in transformer models by incorporating phonetic signals, a modified attention mechanism, and weak supervision guided generation by parts-of-speech constraints.
- CoMix improves performance across four code-mixed tasks: machine translation, sequence classification, named entity recognition (NER), and abstractive summarization. It also achieves new SOTA performance for English-Hinglish translation and NER on LINC Leaderboard and provides better generalization on out-of-domain translation.
- We also propose a new family of metrics based on phonetics and demonstrate that the phonetic variant of BLEU correlates better with human judgement than BLEU on code-mixed text.

Conversational Shopping Assistant: Bridging the gap between online and offline purchase experiences

Accepted @ AMLC 2023

GAURAV ARORA, ET AL.

- This work describes how we built conversational shopping assistant for Amazon (India) which was launched in the IN marketplace for Laptop, Smartphone, and TV categories leading to multi-million dollar incremental OPS.
- In this paper we demonstrate how a conversational system can be bootstrapped successfully without the need for large volumes of chat data and evolved to utilize large language models (LLMs).
- This paper was accepted in Amazon's internal ML conference (AMLC) in 2023 which had acceptance rate of 30%.

DARE: Deep Affordability Recommendation Engine for ART events

Accepted @ AMLC 2022

PRINCE JAIN, GAURAV ARORA, MOHAMMED ABDULLA, VIVEK SEMBIUM

- In this work we present a deep learning based solution which recommends the best affordability combination for each smartphone to minimize cost while achieving the ART (Amazon's sale events) event GMS goals.
- This paper was accepted in Amazon's internal ML conference (AMLC) in 2022 as an oral paper which had acceptance rate of <10%.

Inclusive Speech Detection using Pretrained Language Models

Accepted @ EACL 2021 LT-EDI workshop

GAURAV ARORA*, MEGHA SHARMA*

- In this paper we describe our system that ranked first in Hope Speech Detection (HSD) shared task and fourth in Offensive Language Identification (OLI) shared task, both in Tamil language.

Natural Language Toolkit for Indic Languages - iNLTK

Accepted @ EMNLP 2020 NLP-OSS workshop

GAURAV ARORA

 [goru001/inltk](#)

- iNLTK provides out-of-the-box support for Data Augmentation, Textual Similarity, Sentence Embeddings, Word Embeddings, Tokenization and Text Generation in 13 Indic Languages.
- iNLTK has 100,000+ downloads, 700+ stars on GitHub, and has been widely shared and appreciated on Twitter, LinkedIn, Reddit.
- By using pre-trained models and data augmentation from iNLTK, we achieve more than 95% of the previous best performance by using less than 10% of the training data on publicly available classification datasets.

HINT3: Raising the bar for Intent Detection in the Wild

Accepted @ EMNLP 2020 Insights workshop

GAURAV ARORA, CHIRAG JAIN, MANAS CHATURVEDI, KRUPAL MODI

 [hellohaptik/HINT3](#)

- This paper introduces 3 new datasets created from real user queries received by live chatbots, evaluates popular NLU platforms and highlights critical gaps in language understanding.
- Paper introduces novel subset approach for evaluation and discusses trade-off between out-of-domain vs in-scope performance.

- This paper proposes pre-training ULMFiT on synthetically generated code-mixed data, generated by modelling code-mixed data generation as a Markov process
- Model achieved 0.88 weighted F1-score for code-mixed Tamil-English language in Sub-task B and got 2nd rank on the leader-board

Honors & Awards

Mar. 2021	Indian Achievers Award 2020 , from Indian Achiever's Forum (IAF) in the Young Achievers Category for contribution in nation building through iNLTK	India
Mar. 2020	Fast.ai International Fellow , for contributions to Fast.ai forums	Worldwide
Mar. 2019	Fast.ai International Fellow , for contributions to Fast.ai forums	Worldwide
Dec. 2018	Top-17% , Human Protein Atlas Image Classification, Kaggle , for developing Deep Learning model which classified mixed patterns of proteins in microscope images. The competition had 2172 teams, but I participated individually and hence had 100% contribution in the 366 th placed solution	Worldwide
Oct. 2017	1st Prize, IEEE-Hackathon , for developing chat-bot to help people with emotional decisions in life	PEC
Feb. 2016	Top-100 , among 500,000 students in IT-Olympiad, 2016.	India
Oct. 2016	2nd-Prize, IEEE-Hackathon , for developing an Augmented reality application to help teachers	PEC
Mar. 2016	All India Rank-6 , in IEEE Programming League, among over 1200 undergraduate students.	India
Mar. 2016	2nd Rank , CodeWars, a competitive-programming event hosted by IEEE, PEC on CodeChef	PEC
Nov. 2016	Research Scholarship , for Personal Emotional Doctor - Bot	PEC
2014	All India Rank-885 , JEE-Mains, among 1.4 million candidates	India
2014	1st Rank-Opener, PEC , for best JEE-Mains rank among 600 students of the session 2014-2018.	PEC
2014	1 Lakh Scholarship , from CBSE for 96.4% marks in 12th Boards and 10 CGPA in 10th.	India
2014	Letter of Appreciation , from HRD Ministry, Govt. of India for 96.4% in CBSE-12th exams	India
2011	Catch Them Young , Was among the top-40 students selected from tricity by INFOSYS for 2-week Programming-Basics training on their campus	INFOSYS

Skills & Courses

Mathematics	Discrete Structures for Computer Science, Vector Calculus, Fourier Series and Laplace Transform, Operations Research, Mathematics for Machine Learning: Linear Algebra and Multivariate Calculus
Computer Science	Data Structures and Algorithms, Computer Architecture and Organization, OOP, Microprocessor, DBMS, Operating Systems, Computer Networks, Theory of Computation, Artificial Intelligence, Computer Graphics, Mobile Computing, Fastai: Part 1 and Part 2, DeepLearning.ai by Andrew Ng, Deep Learning by Prof. Mitesh Khapra, IIT Madras
Programming & Web	Python, C, C++, Javascript, TypeScript, EcmaScript6, AngularJS, ReactJS, Angular4, Webpack, Django with Python
Machine Learning tools	Pytorch, Pandas, Numpy, ScikitLearn, SciPy, Fastai, Transformers library

Extracurricular Activity

- Given more than 10 motivational lectures at schools, college etc to an audience of about 100, talking about *student-life challenges* and *success-hacks*.